

# House Sales in King County, USA

This dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015.

Variable	Description
id	A notation for a house
date	Date house was sold
price	Price is prediction target
bedrooms	Number of bedrooms
bathrooms	Number of bathrooms
sqft_living	Square footage of the home
sqft_lot	Square footage of the lot
floors	Total floors (levels) in house
waterfront	House which has a view to a waterfront
view	Has been viewed
condition	How good the condition is overall
grade	overall grade given to the housing unit, based on King County grading system
sqft_above	Square footage of house apart from basement
sqft_basement	Square footage of the basement
yr_built	Built Year
yr_renovated	Year when house was renovated
zipcode	Zip code
lat	Latitude coordinate
long	Longitude coordinate
sqft_living15	Living room area in 2015(implies-- some renovations) This might or might not have affected the lotsize area
sqft_lot15	LotSize area in 2015(implies-- some renovations)

```

#After executing the below command restart the kernel and run all
cells.
!pip3 install scikit-learn --upgrade --user

Requirement already satisfied: scikit-learn in c:\users\siranjeevi\
appdata\local\programs\python\python311\lib\site-packages (1.3.1)
Collecting scikit-learn
  Downloading scikit_learn-1.3.2-cp311-cp311-win_amd64.whl.metadata
(11 kB)
Requirement already satisfied: numpy<2.0,>=1.17.3 in c:\users\
siranjeevi\appdata\local\programs\python\python311\lib\site-packages
(from scikit-learn) (1.25.2)
Requirement already satisfied: scipy>=1.5.0 in c:\users\siranjeevi\
appdata\local\programs\python\python311\lib\site-packages (from
scikit-learn) (1.11.2)
Requirement already satisfied: joblib>=1.1.1 in c:\users\siranjeevi\
appdata\local\programs\python\python311\lib\site-packages (from
scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\
siranjeevi\appdata\local\programs\python\python311\lib\site-packages
(from scikit-learn) (3.2.0)
Downloading scikit_learn-1.3.2-cp311-cp311-win_amd64.whl (9.2 MB)
----- 0.0/9.2 MB ? eta -:--:--
-- 0.7/9.2 MB 14.2 MB/s eta
0:00:01
----- 2.2/9.2 MB 23.1 MB/s eta
0:00:01
----- 4.3/9.2 MB 30.4 MB/s eta
0:00:01
----- 5.2/9.2 MB 27.8 MB/s eta
0:00:01
----- 6.2/9.2 MB 26.3 MB/s eta
0:00:01
----- 7.0/9.2 MB 25.0 MB/s eta
0:00:01
----- 7.9/9.2 MB 24.2 MB/s eta
0:00:01
----- 8.8/9.2 MB 23.5 MB/s eta
0:00:01
----- 9.2/9.2 MB 23.6 MB/s eta
0:00:01
----- 9.2/9.2 MB 21.8 MB/s eta
0:00:00
Installing collected packages: scikit-learn
Successfully installed scikit-learn-1.3.2

```

You will require the following libraries:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression
%matplotlib inline
```

## Module 1: Importing Data Sets

Load the csv:

```
file_name='https://cf-courses-data.s3.us.cloud-object-
storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-
SkillsNetwork/labs/FinalModule_Coursera/data/kc_house_data_NaN.csv'
df=pd.read_csv(file_name)
```

We use the method head to display the first 5 columns of the dataframe.

```
df.head()
```

	Unnamed: 0	id	date	price	bedrooms
0	0	7129300520	20141013T000000	221900.0	3.0
1	1	6414100192	20141209T000000	538000.0	3.0
2	2	5631500400	20150225T000000	180000.0	2.0
3	3	2487200875	20141209T000000	604000.0	4.0
4	4	1954400510	20150218T000000	510000.0	3.0

  

	sqft_living	sqft_lot	floors	waterfront	...	grade
0	1180	5650	1.0	0	...	7
1	2570	7242	2.0	0	...	7
2	770	10000	1.0	0	...	6
3	1960	5000	1.0	0	...	7
4	1680	8080	1.0	0	...	8

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0	1955	0	98178	47.5112	-122.257	
1	400	1951	1991	98125	47.7210	-122.319	
2	0	1933	0	98028	47.7379	-122.233	
3	910	1965	0	98136	47.5208	-122.393	
4	0	1987	0	98074	47.6168	-122.045	

  

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 22 columns]

## Question 1

Display the data types of each column using the function `dtypes`, then take a screenshot and submit it, include your code in the image.

```
# Display the data types of each column
data_types = df.dtypes
```

```
# Print the result
print(data_types)
```

```
Unnamed: 0      int64
id              int64
date            object
price           float64
bedrooms        float64
bathrooms       float64
sqft_living      int64
sqft_lot         int64
floors           float64
waterfront       int64
view             int64
condition        int64
grade            int64
sqft_above       int64
sqft_basement    int64
yr_built         int64
yr_renovated     int64
zipcode          int64
lat              float64
long             float64
sqft_living15    int64
```

```
sqft_lot15      int64
dtype: object
```

We use the method describe to obtain a statistical summary of the dataframe.

```
df.describe()
```

```
      Unnamed: 0      id      price      bedrooms
bathrooms \
count  21613.000000  2.161300e+04  2.161300e+04  21600.000000
21603.000000
mean   10806.000000  4.580302e+09  5.400881e+05      3.372870
2.115736
std    6239.28002   2.876566e+09  3.671272e+05      0.926657
0.768996
min      0.000000   1.000102e+06  7.500000e+04      1.000000
0.500000
25%     5403.000000  2.123049e+09  3.219500e+05      3.000000
1.750000
50%     10806.000000  3.904930e+09  4.500000e+05      3.000000
2.250000
75%     16209.000000  7.308900e+09  6.450000e+05      4.000000
2.500000
max     21612.000000  9.900000e+09  7.700000e+06     33.000000
8.000000

      sqft_living      sqft_lot      floors      waterfront
view \
count  21613.000000  2.161300e+04  21613.000000  21613.000000
21613.000000
mean    2079.899736  1.510697e+04      1.494309      0.007542
0.234303
std      918.440897  4.142051e+04      0.539989      0.086517
0.766318
min      290.000000  5.200000e+02      1.000000      0.000000
0.000000
25%     1427.000000  5.040000e+03      1.000000      0.000000
0.000000
50%     1910.000000  7.618000e+03      1.500000      0.000000
0.000000
75%     2550.000000  1.068800e+04      2.000000      0.000000
0.000000
max     13540.000000  1.651359e+06      3.500000      1.000000
4.000000

      ...      grade      sqft_above      sqft_basement      yr_built \
count  ...  21613.000000  21613.000000  21613.000000  21613.000000
mean   ...      7.656873   1788.390691     291.509045   1971.005136
std    ...      1.175459    828.090978    442.575043    29.373411
```

min	...	1.000000	290.000000	0.000000	1900.000000
25%	...	7.000000	1190.000000	0.000000	1951.000000
50%	...	7.000000	1560.000000	0.000000	1975.000000
75%	...	8.000000	2210.000000	560.000000	1997.000000
max	...	13.000000	9410.000000	4820.000000	2015.000000

	yr_renovated	zipcode	lat	long
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	84.402258	98077.939805	47.560053	-122.213896
std	401.679240	53.505026	0.138564	0.140828
min	0.000000	98001.000000	47.155900	-122.519000
25%	0.000000	98033.000000	47.471000	-122.328000
50%	0.000000	98065.000000	47.571800	-122.230000
75%	0.000000	98118.000000	47.678000	-122.125000
max	2015.000000	98199.000000	47.777600	-121.315000

	sqft_lot15
count	21613.000000
mean	12768.455652
std	27304.179631
min	651.000000
25%	5100.000000
50%	7620.000000
75%	10083.000000
max	871200.000000

[8 rows x 21 columns]

## Module 2: Data Wrangling

### Question 2

Drop the columns "id" and "Unnamed: 0" from axis 1 using the method drop(), then use the method describe() to obtain a statistical summary of the data. Take a screenshot and submit it, make sure the inplace parameter is set to True

```
# Drop the specified columns from axis 1
df.drop(["id", "Unnamed: 0"], axis=1, inplace=True)
```

```
# Use describe() to obtain a statistical summary of the data
summary = df.describe()
```

```
# Print the summary
print(summary)
```

	price	bedrooms	bathrooms	sqft_living
count	2.161300e+04	21613.000000	21613.000000	21613.000000
mean	5.400881e+05	3.372870	2.115736	2079.899736
std	3.671272e+05	0.926378	0.768818	918.440897
min	7.500000e+04	1.000000	0.500000	290.000000
25%	3.219500e+05	3.000000	1.750000	1427.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000
75%	6.450000e+05	4.000000	2.500000	2550.000000
max	7.700000e+06	33.000000	8.000000	13540.000000

	floors	waterfront	view	condition
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.494309	0.007542	0.234303	3.409430
std	0.539989	0.086517	0.766318	0.650743
min	1.000000	0.000000	0.000000	1.000000
25%	1.000000	0.000000	0.000000	3.000000
50%	1.500000	0.000000	0.000000	3.000000
75%	2.000000	0.000000	0.000000	4.000000
max	3.500000	1.000000	4.000000	5.000000

	sqft_above	sqft_basement	yr_built	yr_renovated
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	1788.390691	291.509045	1971.005136	84.402258

98077.939805				
std	828.090978	442.575043	29.373411	401.679240
53.505026				
min	290.000000	0.000000	1900.000000	0.000000
98001.000000				
25%	1190.000000	0.000000	1951.000000	0.000000
98033.000000				
50%	1560.000000	0.000000	1975.000000	0.000000
98065.000000				
75%	2210.000000	560.000000	1997.000000	0.000000
98118.000000				
max	9410.000000	4820.000000	2015.000000	2015.000000
98199.000000				

  

	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	47.560053	-122.213896	1986.552492	12768.455652
std	0.138564	0.140828	685.391304	27304.179631
min	47.155900	-122.519000	399.000000	651.000000
25%	47.471000	-122.328000	1490.000000	5100.000000
50%	47.571800	-122.230000	1840.000000	7620.000000
75%	47.678000	-122.125000	2360.000000	10083.000000
max	47.777600	-121.315000	6210.000000	871200.000000

We can see we have missing values for the columns bedrooms and bathrooms

```
print("number of NaN values for the column bedrooms :",
df['bedrooms'].isnull().sum())
print("number of NaN values for the column bathrooms :",
df['bathrooms'].isnull().sum())
```

```
number of NaN values for the column bedrooms : 13
number of NaN values for the column bathrooms : 10
```

We can replace the missing values of the column 'bedrooms' with the mean of the column 'bedrooms' using the method replace(). Don't forget to set the inplace parameter to True

```
mean=df['bedrooms'].mean()
df['bedrooms'].replace(np.nan,mean, inplace=True)
```

We also replace the missing values of the column 'bathrooms' with the mean of the column 'bathrooms' using the method replace(). Don't forget to set the inplace parameter to True

```
mean=df['bathrooms'].mean()
df['bathrooms'].replace(np.nan,mean, inplace=True)

print("number of NaN values for the column bedrooms :",
df['bedrooms'].isnull().sum())
```



```
print("number of NaN values for the column bathrooms :",  
df['bathrooms'].isnull().sum())
```

```
number of NaN values for the column bedrooms : 0  
number of NaN values for the column bathrooms : 0
```

## Module 3: Exploratory Data Analysis

### Question 3

Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a dataframe.

```
# Count the number of houses with unique floor values  
floor_counts = df['floors'].value_counts()  
  
# Convert the result to a DataFrame using to_frame()  
floor_counts_df = floor_counts.to_frame()  
  
# Print the DataFrame  
print(floor_counts_df)
```

	count
1.0	10680
2.0	8241
1.5	1910
3.0	613
2.5	161
3.5	8

### Question 4

Use the function `boxplot` in the `seaborn` library to determine whether houses with a waterfront view or without a waterfront view have more price outliers.

```
import seaborn as sns  
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(12, 6))  
sns.boxplot(x='waterfront', y='price', data=df)  
plt.title('Boxplot of Price for Houses with and without Waterfront View')  
plt.show()
```



## Question 5

Use the function `regplot` in the `seaborn` library to determine if the feature `sqft_above` is negatively or positively correlated with price.

```
plt.figure(figsize=(12, 6))
sns.regplot(x='sqft_above', y='price', data=df)
plt.title('Regplot of Price vs. sqft_above')
plt.show()
```



We can use the Pandas method `corr()` to find the feature other than price that is most correlated with price.

```
df.corr()['price'].sort_values()
```

```
-----
-----
ValueError                                Traceback (most recent call
last)
Cell In[22], line 2
      1 # Use the Pandas method corr() to find the correlation between
features and price
----> 2 correlation_matrix = df.corr()
      4 # Extract the correlation of features with the target variable
'price'
      5 correlation_with_price = correlation_matrix['price']
```

```
File c:\Users\siranjeevi\AppData\Local\Programs\Python\Python311\Lib\
site-packages\pandas\core\frame.py:10704, in DataFrame.corr(self,
method, min_periods, numeric_only)
    10702 cols = data.columns
    10703 idx = cols.copy()
> 10704 mat = data.to_numpy(dtype=float, na_value=np.nan, copy=False)
    10706 if method == "pearson":
    10707     correl = libalgos.nancorr(mat, minp=min_periods)
```

```
File c:\Users\siranjeevi\AppData\Local\Programs\Python\Python311\Lib\
site-packages\pandas\core\frame.py:1889, in DataFrame.to_numpy(self,
dtype, copy, na_value)
```

```

1887 if dtype is not None:
1888     dtype = np.dtype(dtype)
-> 1889 result = self._mgr.as_array(dtype=dtype, copy=copy,
na_value=na_value)
1890 if result.dtype is not dtype:
1891     result = np.array(result, dtype=dtype, copy=False)

File c:\Users\siranjeevi\AppData\Local\Programs\Python\Python311\Lib\
site-packages\pandas\core\internals\managers.py:1656, in
BlockManager.as_array(self, dtype, copy, na_value)
1654     arr.flags.writeable = False
1655 else:
-> 1656     arr = self._interleave(dtype=dtype, na_value=na_value)
1657     # The underlying data was copied within _interleave, so no
need
1658     # to further copy if copy=True or setting na_value
1660 if na_value is lib.no_default:

File c:\Users\siranjeevi\AppData\Local\Programs\Python\Python311\Lib\
site-packages\pandas\core\internals\managers.py:1715, in
BlockManager._interleave(self, dtype, na_value)
1713     else:
1714         arr = blk.get_values(dtype)
-> 1715     result[rl.indexer] = arr
1716     itemmask[rl.indexer] = 1
1718 if not itemmask.all():

ValueError: could not convert string to float: '20141013T000000'

```

## Module 4: Model Development

We can Fit a linear regression model using the longitude feature 'long' and caculate the R<sup>2</sup>.

```

X = df[['long']]
Y = df['price']
lm = LinearRegression()
lm.fit(X,Y)
lm.score(X, Y)

0.00046769430149007363

```

### Question 6

Fit a linear regression model to predict the 'price' using the feature 'sqft\_living' then calculate the R<sup>2</sup>. Take a screenshot of your code and the value of the R<sup>2</sup>.

```

from sklearn.linear_model import LinearRegression

# Feature and target variable
X_sqft_living = df[['sqft_living']]
Y_price = df['price']

# Create a linear regression model
lm_sqft_living = LinearRegression()

# Fit the model
lm_sqft_living.fit(X_sqft_living, Y_price)

# Calculate R^2
r_squared_sqft_living = lm_sqft_living.score(X_sqft_living, Y_price)

# Print the R^2 value
print(f"R^2 for linear regression using 'sqft_living': {r_squared_sqft_living}")

R^2 for linear regression using 'sqft_living': 0.4928532179037931

```

## Question 7

Fit a linear regression model to predict the 'price' using the list of features:

```

features = ["floors",
"waterfront", "lat" , "bedrooms" , "sqft_basement" , "view" , "bathrooms", "
sqft_living15", "sqft_above", "grade", "sqft_living"]

```

Then calculate the R<sup>2</sup>. Take a screenshot of your code.

```

from sklearn.linear_model import LinearRegression

# List of features
features = ["floors", "waterfront", "lat", "bedrooms",
"sqft_basement", "view", "bathrooms",
            "sqft_living15", "sqft_above", "grade", "sqft_living"]

# Feature and target variable
X_features = df[features]
Y_price = df['price']

# Create a linear regression model
lm_features = LinearRegression()

# Fit the model
lm_features.fit(X_features, Y_price)

# Calculate R^2
r_squared_features = lm_features.score(X_features, Y_price)

```

```
# Print the R^2 value
print(f"R^2 for linear regression using the given features:
{r_squared_features}")
```

R^2 for linear regression using the given features: 0.6576372970735713

## This will help with Question 8

Create a list of tuples, the first element in the tuple contains the name of the estimator:

'scale'

'polynomial'

'model'

The second element in the tuple contains the model constructor

StandardScaler()

PolynomialFeatures(include\_bias=False)

LinearRegression()

```
Input=[('scale',StandardScaler()),('polynomial',
PolynomialFeatures(include_bias=False)),('model',LinearRegression())]
```

## Question 8

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list features, and calculate the R^2.

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression

# List of tuples for the pipeline
pipeline_steps = [
    ('scale', StandardScaler()),
    ('polynomial', PolynomialFeatures(include_bias=False)),
    ('model', LinearRegression())
]

# Create a pipeline object
pipeline = Pipeline(pipeline_steps)

# Fit the pipeline using the specified features
pipeline.fit(df[features], df['price'])

# Calculate R^2
```

```
r_squared_pipeline = pipeline.score(df[features], df['price'])

# Print the R^2 value
print(f"R^2 for the pipeline using the given features:
{r_squared_pipeline}")

R^2 for the pipeline using the given features: 0.7508598253545078
```

## Module 5: Model Evaluation and Refinement

Import the necessary modules:

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
print("done")

done
```

We will split the data into training and testing sets:

```
features = ["floors",
"waterfront", "lat" , "bedrooms" , "sqft_basement" , "view" , "bathrooms", "
sqft_living15", "sqft_above", "grade", "sqft_living"]
X = df[features]
Y = df['price']

x_train, x_test, y_train, y_test = train_test_split(X, Y,
test_size=0.15, random_state=1)

print("number of test samples:", x_test.shape[0])
print("number of training samples:", x_train.shape[0])

number of test samples: 3242
number of training samples: 18371
```

### Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R<sup>2</sup> using the test data.

```
from sklearn.linear_model import Ridge

from sklearn.linear_model import Ridge
from sklearn.model_selection import train_test_split

# Features and target variable
X = df[features]
```

```

Y = df['price']

# Split the data into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(X, Y,
test_size=0.15, random_state=1)

# Create and fit a Ridge regression object with regularization
parameter set to 0.1
ridge_model = Ridge(alpha=0.1)
ridge_model.fit(x_train, y_train)

# Calculate R^2 using the test data
r_squared_ridge = ridge_model.score(x_test, y_test)

# Print the R^2 value
print(f"R^2 for Ridge regression with regularization parameter 0.1:
{r_squared_ridge}")

R^2 for Ridge regression with regularization parameter 0.1:
0.6478759163939116

```

## Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the  $R^2$  utilising the test data provided. Take a screenshot of your code and the  $R^2$ .

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline

# Perform a second order polynomial transform on both training and
testing data
degree = 2
poly = PolynomialFeatures(degree=degree)

x_train_poly = poly.fit_transform(x_train)
x_test_poly = poly.transform(x_test)

# Create and fit a Ridge regression object with regularization
parameter set to 0.1
ridge_model_poly = Ridge(alpha=0.1)
ridge_model_poly.fit(x_train_poly, y_train)

# Calculate R^2 using the test data
r_squared_poly = ridge_model_poly.score(x_test_poly, y_test)

# Print the R^2 value
print(f"R^2 for Ridge regression with polynomial transform (degree
{degree}) and regularization parameter 0.1: {r_squared_poly}")

```



R<sup>2</sup> for Ridge regression with polynomial transform (degree 2) and regularization parameter 0.1: 0.700274426566343

Joseph Santarcangelo has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: Michelle Carey, Mavis Zhou

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2022-07-29	2.3	Lakshmi Holla	Added library import
2020-12-01	2.2	Aije Egwaikhide	Converted Data description from text to table
2020-10-06	2.1	Lakshmi Holla	Changed markdown instruction of Question1
2020-08-27	2.0	Malika Singla	Added lab to GitLab

© IBM Corporation 2020. All rights reserved.